

# Week 09 - Lab Session Results

February 26, 2025

## Text Representation

### Exercise 1

Load the [metadata file](#) from Absalon and discard any item that was not rated by our subset of users (not in training or test sets). You can refer to the [original metadata file](#) if you want to look up more explanations about the columns of the metadata file. Apply preprocessing in this order: lowercasing, tokenizing, stemming, and stopwords removal (including punctuation) to clean up the text from the title. Report the vocabulary size before and after the preprocessing. You may have to specify the language for these steps.

Total number of items: 84

Vocabulary size before preprocessing: 545

Vocabulary size after preprocessing: 471

### Exercise 2

Representation in vector spaces.

#### 2.1

Represent all the items from Exercise 1 in a TF-IDF space. Interpret the meaning of the TF-IDF matrix dimensions. Be careful with multiple instances of preprocessing in the process, as default settings for creating the TF-IDF space may include some.

Tip: You may use the library [scikit-learn](#)

TF-IDF matrix shape: (84, 471)

#### 2.2

Using the TF-IDF representation, compute and the cosine similarity between products with asin B000FI4S1E, B000LIBUBY and B000W0C07Y. Take a look at their features to see whether results make sense with their characteristics. Round your final answer to 3 decimal places.

Similarity between 'B000FI4S1E' and 'B000LIBUBY': 0.031

Similarity between 'B000FI4S1E' and 'B000W0C07Y': 0.024

Similarity between 'B000LIBUBY' and 'B000W0C07Y': 0.501

## Exercise 3

Representation in vector spaces with contextual Word Embeddings.

### 3.1.

Represent all the products from Exercise 1 in a vector space using embeddings from a pre-trained BERT model. The final embedding of a product should be the average of the word embeddings from all the words in the 'title'. Critically evaluate this procedure.

What is the vocabulary size of the model? What are the dimensions of the last hidden state?

Tip: you may install the transformers library and use their pretrained [BERT model uncased](#).

Using device: cpu

Vocabulary size of 30522. Input dimension: 768.

last\_hidden\_states: torch.Size([84, 52, 768])

### 3.2.

Using the representation obtained from Exercise 3.1., compute the cosine similarity between items with asin B000FI4S1E, B000LIBUBY and B000W0C07Y. Round your final answer to 3 decimal places.

Similarity between 'B000FI4S1E' and 'B000LIBUBY': 0.734

Similarity between 'B000FI4S1E' and 'B000W0C07Y': 0.659

Similarity between 'B000LIBUBY' and 'B000W0C07Y': 0.748